

SF2935: Modern Methods of Statistical Learning (Autumn 2025)

This advanced course in statistical learning dives into modern methods and their theoretical foundations. While the core focus is on conceptual understanding and statistical principles, the material will be reinforced through empirical demonstrations using data.

The basic course syllabus is available at <https://www.kth.se/student/kurser/kurs/SF2935/?l=en>. This document provides a detailed course plan, including information on the main texts, tentative schedule, assessment structure, and recommended supplementary resources.

1 Prerequisites

Students are expected to have a strong¹ background in linear algebra, calculus, probability and statistics, and proficiency in a programming language suitable for scientific computing. A prior introductory course in machine learning or statistical modeling is highly recommended.

2 Course Materials

The main text for this course is:

- ***The Elements of Statistical Learning: Data Mining, Inference, and Prediction*** (ESL), by Trevor Hastie, Robert Tibshirani, and Jerome Friedman (2nd edition, corrected 12th printing, 2017).
Available at: <https://hastie.su.domains/ElemStatLearn/>

ESL is widely regarded as the gold standard reference for statistical learning, offering comprehensive coverage of both theory and methodology. We aim to cover approximately 70% of the materials in ESL, supplemented with:

- ***An Introduction to Statistical Learning: With Applications in Python*** (ISL), by Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani (2023).
Available at: <https://www.statlearning.com/>
- ***Learning Theory From First Principles*** (Bach), by Francis Bach (final version, 2025).
Available at: https://www.di.ens.fr/~fbach/lftp_book.pdf

ISL emphasizes a conceptual and application-oriented understanding of statistical learning methods, while Bach provides a rigorous treatment of the underlying mathematical principles. Additional materials may also be used to supplement the main text.

Lecture slides and the associated code will be made available after each class.

3 Tentative Schedule

We have 14 lecture sessions and 7 tutorial sessions. All the sessions are held from 13:00-15:00 except the sessions on Aug 29 and Sep 19 (15:00-17:00).

¹At a minimum, you should be familiar with most of the material in Part I of: <https://mml-book.github.io/>.

Date	Session	Topic	References
Aug 26	L1	Introduction to Supervised Learning	ESL (Ch. 1–2), ISL (Ch. 1–2), Bach (Ch. 2)
Aug 28	L2	Linear Methods for Regression	ESL (Ch. 3), ISL (Ch. 3, Ch. 6), Bach (Ch. 3)
Aug 29*	T1	Tutorial	-
Sep 2	L3	Linear Methods for Classification	ESL (Ch. 4), ISL (Ch. 4)
Sep 4	L4	Maximal Margin Methods for Classification	ESL (Ch. 4, Ch. 12.1-12.3), ISL (Ch. 9), Bach (Ch. 4.1)
Sep 5	T2	Tutorial	-
Sep 9	L5	Beyond Linearity: Basis Expansions, Local Averaging, Kernel Smoothing	ESL (Ch. 5.1-5.7, Ch. 6), ISL (Ch. 7), Bach (Ch. 6)
Sep 11	L6	Kernel Methods: Fundamentals	Bach (Ch. 7), ESL (Ch. 5.8)
Sep 12	T3	Tutorial, Project Proposal Due	-
Sep 16	L7	Kernel Methods: Special Topics	Bach (Ch. 7), additional materials
Sep 18	L8	Model Assessment and Selection, Generalization Theory	ESL (Ch. 7, 9.1), ISL (Ch. 5), Bach (Ch. 4)
Sep 19*	T4	Tutorial	-
Sep 23	L9	Tree-Based Methods, Ensembling Methods	ESL (Ch. 9.2-10, Ch. 15), ISL (Ch. 8), Bach (Ch. 10)
Sep 25	L10	Neural Networks: Fundamentals	ISL (Ch. 10), ESL (Ch. 11), Bach (Ch. 9)
Sep 26	15	Tutorial	-
Sep 30	L11	Rethinking Generalization: Modern Phenomena in Overparameterized Models	ISL (Ch. 10), ESL (Ch. 11), Bach (Ch. 12)
Oct 2	L12	Rethinking Robustness: Probabilistic Approaches to Reliable Models	ESL (Ch. 8, 11), Bach (Ch. 14), additional materials
Oct 3	T6	Tutorial	-
Oct 7	L13	Unsupervised Learning: Fundamentals	ESL (Ch. 14), ISL (Ch. 12)
Oct 9	L14	Unsupervised Learning: Special Topics	Additional materials
Oct 10	T7	Tutorial	-
Oct 17	-	Project Report Due	-
Oct 23	-	Exam	-
Dec 18	-	Re-Exam	-

4 Assessments and Grading

For the “*assignments*” (3.0 credits), the grading scale is P or F (P is awarded as long as a complete assignment is submitted by the deadline). For the “*examination*” (4.5 credits) and the final course grade (A, B, C, D, E, FX, F), it will be based on the following equally weighted² components:

- A **course project**: 50%
- An **exam**: 50%

Please refer to the Lecture 1 slides for tips on how to succeed in the course. A mock exam will be posted during the week of October 7, and the format of the actual exam will be the same. Past exams from previous editions of the course, which were taught by a different instructor, may not be useful since the materials covered are not exactly the same. For the exam, you are allowed to bring one double-sided A4 cheatsheet.

5 Additional Resources

The following references may also be useful for deepening your understanding of the topics covered:

- *Machine Learning: A First Course for Engineers and Scientists*, by Andreas Lindholm, Niklas Wahlström, Fredrik Lindsten, Thomas Schön (2022)
- *Pattern Recognition and Machine Learning*, by Christopher Bishop (2006)
- *Deep Learning: Foundations and Concepts*, by Christopher Bishop and Hugh Bishop (2024)
- *Probabilistic Machine Learning: An Introduction*, by Kevin Murphy (2022)
- *Patterns, Predictions, and Actions*, by Moritz Hardt and Benjamin Recht (2022)
- *Understanding Machine Learning: From Theory to Algorithms*, by Shai Shalev-Shwartz and Shai Ben-David (2014)
- *Foundations of Machine Learning*, by Mehryar Mohri, Afshin Rostamizadeh, Ameet Talwalkar (2nd edition, 2018)
- *Learning from Data*, by Yaser Abu-Mostafa, Malik Magdon-Ismael, Hsuan-Tien Lin (2012)
- *Deep Learning*, by Ian Goodfellow, Yoshua Bengio, Aaron Courville (2016)
- *Mathematics for Machine Learning*, by Marc Deisenroth, Aldo Faisal, Cheng Soon Ong (2020)
- *Scikit-learn: Machine Learning in Python*^{*} (available at <https://scikit-learn.org/stable/>)
- *Dive into Deep Learning (D2L)*^{*}, by Aston Zhang, Zachary Lipton, Mu Li, Alexander Smola (available at <https://d2l.ai/>)
- *Practical Deep Learning for Coders*^{*}, by Jeremy Howard (at <https://course.fast.ai/>)
- *PyTorch in One Hour: From Tensors to Training Neural Networks on Multiple GPUs*^{*}, by Sebastian Raschka (available at <https://sebastianraschka.com/teaching/pytorch-1h/>)

^{*}: useful references for coding and implementation in Python and PyTorch

²The project and exam are each worth up to 50% of the final grade. The project can only boost your total score — if the project score is lower than the exam score, the exam score will determine the final outcome.